

Heart Disease-

January 20, 2024

0.1 Import Libraries

```
[2]: import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns
      import sklearn as sk
```

0.2 Load the dataset

```
[3]: df=pd.read_csv("Heart.csv")
```

0.3 Analyze the dataset

```
[4]: df.head()
```

```
[4]:   Age Sex ChestPainType RestingBP Cholesterol FastingBS RestingECG MaxHR \
0    40   M           ATA     140.0      289.0        0.0    Normal  172.0
1    49   F           NAP     160.0      180.0        0.0    Normal  156.0
2    37   M           ATA     130.0      283.0        0.0       ST  98.0
3    48   F           ASY      NaN      214.0        0.0    Normal  108.0
4    54   M           NAP     150.0      195.0        0.0    Normal  122.0

   ExerciseAngina Oldpeak ST_Slope HeartDisease
0              N     0.0      Up          0
1              N     1.0     Flat         1
2              N     0.0      Up          0
3              Y     1.5     Flat         1
4              N     0.0      Up          0
```

```
[5]: df.tail()
```

```
[5]:   Age Sex ChestPainType RestingBP Cholesterol FastingBS RestingECG \
913   45   M           TA     110.0      264.0        0.0    Normal
914   68   M           ASY     144.0      193.0        1.0    Normal
915   57   M           ASY     130.0      131.0        0.0    Normal
916   57   F           ATA     130.0      236.0        0.0       LVH
917   38   M           NAP     138.0      175.0        0.0    Normal
```

	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
913	132.0	N	1.2	Flat	1
914	141.0	N	3.4	Flat	1
915	115.0	Y	1.2	Flat	1
916	174.0	N	0.0	Flat	1
917	173.0	N	0.0	Up	0

0.4 Describe the dataset

[6]: df.describe()

```
[6]:          Age   RestingBP Cholesterol   FastingBS      MaxHR \
count  918.000000  917.000000  917.000000  917.000000  917.000000
mean   53.510893  132.390403  199.016358  0.233370  136.801527
std    9.432617  18.523331  109.246330  0.423206  25.473119
min   28.000000  0.000000  0.000000  0.000000  60.000000
25%  47.000000  120.000000  174.000000  0.000000  120.000000
50%  54.000000  130.000000  223.000000  0.000000  138.000000
75%  60.000000  140.000000  267.000000  0.000000  156.000000
max  77.000000  200.000000  603.000000  1.000000  202.000000

          Oldpeak HeartDisease
count  918.000000  918.000000
mean   0.887364  0.553377
std    1.066570  0.497414
min   -2.600000  0.000000
25%  0.000000  0.000000
50%  0.600000  1.000000
75%  1.500000  1.000000
max  6.200000  1.000000
```

0.5 Checking which columns contain null values

[7]: df.isnull().sum()

```
[7]: Age        0
Sex         0
ChestPainType 0
RestingBP    1
Cholesterol  1
FastingBS    1
RestingECG   0
MaxHR        1
ExerciseAngina 0
Oldpeak      0
ST_Slope     0
```

```
HeartDisease      0  
dtype: int64
```

0.6 info()

```
[8]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 918 entries, 0 to 917  
Data columns (total 12 columns):  
 #   Column            Non-Null Count  Dtype     
---  --     
 0   Age               918 non-null    int64    
 1   Sex               918 non-null    object    
 2   ChestPainType    918 non-null    object    
 3   RestingBP         917 non-null    float64   
 4   Cholesterol      917 non-null    float64   
 5   FastingBS         917 non-null    float64   
 6   RestingECG        918 non-null    object    
 7   MaxHR             917 non-null    float64   
 8   ExerciseAngina   918 non-null    object    
 9   Oldpeak           918 non-null    float64   
 10  ST_Slope          918 non-null    object    
 11  HeartDisease     918 non-null    int64    
dtypes: float64(5), int64(2), object(5)  
memory usage: 86.2+ KB
```

```
[9]: df.shape
```

```
[9]: (918, 12)
```


```
[10]: sns.countplot(data=df,x='Sex');  
plt.title('Sex Distribution')  
plt.xlabel('Sex')  
plt.ylabel('Count')
```

```
[10]: Text(0, 0.5, 'Count')
```




```
[11]: sns.displot(df['Age'])
```

```
[11]: <seaborn.axisgrid.FacetGrid at 0x1c75f9fdc90>
```




```
[12]: sns.displot(df['RestingBP'])
```

```
[12]: <seaborn.axisgrid.FacetGrid at 0x1c75f9e3310>
```



```
[13]: sns.displot(df['Cholesterol'])
```

```
[13]: <seaborn.axisgrid.FacetGrid at 0x1c765848c10>
```



0.7 Duplicate Values

```
[14]: df_data=df.duplicated().any()
```

```
[15]: df_data
```

```
[15]: False
```

0.7.1 label

```
[16]: df.HeartDisease.value_counts()
```

```
[16]: 1      508  
0      410  
Name: HeartDisease, dtype: int64
```

1 completing missing values & outlier

```
[17]: from sklearn.impute import SimpleImputer
numImputer=SimpleImputer(missing_values=np.nan, strategy="most_frequent")
numImputer=numImputer.fit(df[["RestingBP","Cholesterol","FastingBS","MaxHR"]])
new_df=numImputer.transform(df[["RestingBP","Cholesterol","FastingBS","MaxHR"]])
new_df
```

```
[17]: array([[140., 289., 0., 172.],
       [160., 180., 0., 156.],
       [130., 283., 0., 98.],
       ...,
       [130., 131., 0., 115.],
       [130., 236., 0., 174.],
       [138., 175., 0., 173.]])
```

```
[18]: df[["RestingBP","Cholesterol","FastingBS","MaxHR"]]=new_df
df
```


```
[18]:    Age Sex ChestPainType RestingBP Cholesterol FastingBS RestingECG \
0     40   M          ATA      140.0      289.0      0.0    Normal
1     49   F          NAP      160.0      180.0      0.0    Normal
2     37   M          ATA      130.0      283.0      0.0      ST
3     48   F          ASY      120.0      214.0      0.0    Normal
4     54   M          NAP      150.0      195.0      0.0    Normal
..   ...
913    45   M          TA      110.0      264.0      0.0    Normal
914    68   M          ASY      144.0      193.0      1.0    Normal
915    57   M          ASY      130.0      131.0      0.0    Normal
916    57   F          ATA      130.0      236.0      0.0      LVH
917    38   M          NAP      138.0      175.0      0.0    Normal

   MaxHR ExerciseAngina Oldpeak ST_Slope HeartDisease
0     172.0              N     0.0      Up         0
1     156.0              N     1.0     Flat        1
2     98.0               N     0.0      Up         0
3    108.0               Y     1.5     Flat        1
4    122.0              N     0.0      Up         0
..   ...
913   132.0              N     1.2     Flat        1
914   141.0              N     3.4     Flat        1
915   115.0               Y     1.2     Flat        1
916   174.0              N     0.0     Flat        1
917   173.0              N     0.0      Up         0
```


[918 rows x 12 columns]

1.0.1 outlier

```
[19]: plt.boxplot(x=[df["Cholesterol"],df["RestingBP"],df["MaxHR"],df["Oldpeak"]])
plt.xticks([1,2,3,4],[ "Cholesterol", "RestingBP", "MaxHR", "Oldpeak"])
plt.show()
```




```
[20]: plt.scatter(df["Cholesterol"],df["Age"])
plt.xlabel("Cholesterol")
plt.ylabel("Age")
plt.show()
```



```
[21]: df = df.drop(df[(df['RestingBP'] == 0)].index)
```

```
[22]: plt.boxplot(df["RestingBP"])
plt.show()
```



```
[23]: df.loc[df['Cholesterol'] == 0, 'Cholesterol'] = np.nan
```

```
[24]: df.isnull().sum()
```

```
[24]: Age          0  
Sex          0  
ChestPainType 0  
RestingBP     0  
Cholesterol   171  
FastingBS     0  
RestingECG    0  
MaxHR         0  
ExerciseAngina 0  
Oldpeak        0  
ST_Slope       0  
HeartDisease   0  
dtype: int64
```

```
[25]: df["Cholesterol"] = df["Cholesterol"].fillna(df["Cholesterol"].median())  
df["Cholesterol"]
```

```
[25]: 0      289.0
       1      180.0
       2      283.0
       3      214.0
       4      195.0
       ...
      913     264.0
      914     193.0
      915     131.0
      916     236.0
      917     175.0
Name: Cholesterol, Length: 917, dtype: float64
```

2 encoder

Age

```
[26]: df["Sex"].value_counts()
```

```
[26]: M      724
       F      193
Name: Sex, dtype: int64
```

```
[27]: df["Sex"]=df["Sex"].replace(["M","F"],[0,1])
```

ChestPainType

```
[28]: df["ChestPainType"].value_counts()
```

```
[28]: ASY    496
       NAP    202
       ATA    173
       TA     46
Name: ChestPainType, dtype: int64
```

```
[29]: df["ChestPainType"]=df["ChestPainType"].
       ↪replace(["ASY","NAP","ATA","TA"],[0,1,2,3])
```

RestingECG

```
[30]: df["RestingECG"].value_counts()
```

```
[30]: Normal   551
       LVH      188
       ST       178
Name: RestingECG, dtype: int64
```

```
[31]: df["RestingECG"]=df["RestingECG"].replace(["Normal","LVH","ST"],[0,1,2])
```

ExerciseAngina

```
[18]: df["ExerciseAngina"].value_counts()
```

```
[18]: N      546  
Y      371  
Name: ExerciseAngina, dtype: int64
```

```
[19]: df["ExerciseAngina"]=df["ExerciseAngina"].replace(["N","Y"],[0,1])
```

ST_Slope

```
[20]: df["ST_Slope"].value_counts()
```

```
[20]: Flat     459  
Up       395  
Down     63  
Name: ST_Slope, dtype: int64
```

```
[21]: df["ST_Slope"]=df["ST_Slope"].replace(["Flat","Up","Down"],[0,1,2])
```

```
[22]: df
```

```
[22]:    Age  Sex  ChestPainType  RestingBP  Cholesterol  FastingBS  RestingECG \
0     40    0            2        140.0      289.0       0.0          0
1     49    1            1        160.0      180.0       0.0          0
2     37    0            2        130.0      283.0       0.0          2
3     48    1            0        120.0      214.0       0.0          0
4     54    0            1        150.0      195.0       0.0          0
..   ...
913    45    0            3        110.0      264.0       0.0          0
914    68    0            0        144.0      193.0       1.0          0
915    57    0            0        130.0      131.0       0.0          0
916    57    1            2        130.0      236.0       0.0          1
917    38    0            1        138.0      175.0       0.0          0
```

```
      MaxHR  ExerciseAngina  Oldpeak  ST_Slope  HeartDisease
0     172.0            0     0.0       1          0
1     156.0            0     1.0       0          1
2     98.0             0     0.0       1          0
3    108.0             1     1.5       0          1
4    122.0            0     0.0       1          0
..   ...
913    132.0            0     1.2       0          1
914    141.0            0     3.4       0          1
915    115.0             1     1.2       0          1
916    174.0            0     0.0       0          1
917    173.0            0     0.0       1          0
```


```
[917 rows x 12 columns]
```

2.0.1 corr

```
[31]: cor=df.corr()
cor
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	\
Age	1.000000	-0.055670	-0.165951	0.263665	0.045556	
Sex	-0.055670	1.000000	0.168437	-0.011538	0.111122	
ChestPainType	-0.165951	0.168437	1.000000	-0.019763	-0.062737	
RestingBP	0.263665	-0.011538	-0.019763	1.000000	0.086641	
Cholesterol	0.045556	0.111122	-0.062737	0.086641	1.000000	
FastingBS	0.198170	-0.120424	-0.116564	0.068410	0.024851	
RestingECG	0.210694	-0.038758	-0.064901	0.115449	0.010679	
MaxHR	-0.381982	0.189516	0.343310	-0.108317	-0.001780	
ExerciseAngina	0.216017	-0.191226	-0.416553	0.151708	0.072751	
Oldpeak	0.258563	-0.105444	-0.245250	0.173587	0.054583	
ST_Slope	-0.093310	0.066355	0.203046	-0.092236	-0.076421	
HeartDisease	0.282012	-0.305118	-0.471846	0.116984	0.076262	
	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	\
Age	0.198170	0.210694	-0.381982	0.216017	0.258563	
Sex	-0.120424	-0.038758	0.189516	-0.191226	-0.105444	
ChestPainType	-0.116564	-0.064901	0.343310	-0.416553	-0.245250	
RestingBP	0.068410	0.115449	-0.108317	0.151708	0.173587	
Cholesterol	0.024851	0.010679	-0.001780	0.072751	0.054583	
FastingBS	1.000000	0.120381	-0.131195	0.059988	0.053062	
RestingECG	0.120381	1.000000	-0.092709	0.097755	0.100452	
MaxHR	-0.131195	-0.092709	1.000000	-0.369673	-0.160445	
ExerciseAngina	0.059988	0.097755	-0.369673	1.000000	0.409494	
Oldpeak	0.053062	0.100452	-0.160445	0.409494	1.000000	
ST_Slope	-0.044113	-0.020166	0.247979	-0.254216	-0.096810	
HeartDisease	0.267994	0.108440	-0.401137	0.495490	0.403638	
	ST_Slope	HeartDisease				
Age	-0.093310	0.282012				
Sex	0.066355	-0.305118				
ChestPainType	0.203046	-0.471846				
RestingBP	-0.092236	0.116984				
Cholesterol	-0.076421	0.076262				
FastingBS	-0.044113	0.267994				
RestingECG	-0.020166	0.108440				
MaxHR	0.247979	-0.401137				
ExerciseAngina	-0.254216	0.495490				
Oldpeak	-0.096810	0.403638				
ST_Slope	1.000000	-0.397260				
HeartDisease	-0.397260	1.000000				

```
[41]: sns.pairplot(df)
plt.show()
```



2.1 Standardization

```
[31]: from sklearn import preprocessing
X=df.iloc[:, :-1]
Y=df.iloc[:, -1]
```

```
[32]: X
```

```
[32]:      Age   Sex ChestPainType   RestingBP   Cholesterol   FastingBS   RestingECG \
0        40     0            2         140.0       289.0          0.0           0
```

```

1    49    1      1    160.0    180.0    0.0    0
2    37    0      2    130.0    283.0    0.0    2
3    48    1      0    120.0    214.0    0.0    0
4    54    0      1    150.0    195.0    0.0    0
...
913   45    0      3    110.0    264.0    0.0    0
914   68    0      0    144.0    193.0    1.0    0
915   57    0      0    130.0    131.0    0.0    0
916   57    1      2    130.0    236.0    0.0    1
917   38    0      1    138.0    175.0    0.0    0

      MaxHR  ExerciseAngina  Oldpeak  ST_Slope
0     172.0            0       0.0       1
1     156.0            0       1.0       0
2     98.0             0       0.0       1
3    108.0            1       1.5       0
4    122.0            0       0.0       1
...
913   132.0            0       1.2       0
914   141.0            0       3.4       0
915   115.0            1       1.2       0
916   174.0            0       0.0       0
917   173.0            0       0.0       1

```

[917 rows x 11 columns]

[33] : Y

```

[33] : 0      0
1      1
2      0
3      1
4      0
...
913   1
914   1
915   1
916   1
917   0
Name: HeartDisease, Length: 917, dtype: int64

```

3 Split data

```
[34] : from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test=train_test_split(X,Y, test_size=0.3 ,random_state=0)
```

```
print("Train set",X_train.shape, y_train.shape)
print("Test set",X_test.shape, y_test.shape)
```

Train set (641, 11) (641,)
Test set (276, 11) (276,)

```
[35]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X_train_s=sc.fit_transform(X_train)
X_test_s=sc.transform(X_test)
```

```
[36]: X_train_s
```

```
[36]: array([[-2.26099369, -0.51894791, -0.82065891, ..., -0.83412661,
           -0.85419772, -0.89779876],
           [-1.84109019, -0.51894791,  0.22932584, ..., -0.83412661,
            0.09359653, -0.89779876],
           [-0.26645203, -0.51894791, -0.82065891, ..., -0.83412661,
            1.23094963, -0.89779876],
           ...,
           [ 1.83306551,  1.92697569, -0.82065891, ..., -0.83412661,
            0.66227308, -0.89779876],
           [ 0.0484756 , -0.51894791,  0.22932584, ..., -0.83412661,
            -0.66463887,  0.70971243],
           [ 0.78330674, -0.51894791, -0.82065891, ...,  1.19885877,
            2.55786158, -0.89779876]])
```

```
[37]: from sklearn.feature_selection import VarianceThreshold
variance_selector = VarianceThreshold(threshold=0)
X_train_fs = variance_selector.fit_transform(X_train_s)
X_test_fs = variance_selector.transform(X_test_s)
print(f"{X_train.shape[1]-X_train_fs.shape[1]} features have been removed,\n"
      f"{X_train_fs.shape[1]} features remain")
```

0 features have been removed, 11 features remain

```
[38]: from sklearn.feature_selection import SelectKBest, f_classif
selector = SelectKBest(f_classif, k='all')
X_train_fs1 = selector.fit_transform(X_train_s, y_train)
X_test_fs1 = selector.transform(X_test_s)
print(f"{X_train.shape[1]-X_train_fs1.shape[1]} features have been removed,\n"
      f"{X_train_fs1.shape[1]} features remain")
```

0 features have been removed, 11 features remain

4 KNN

```
[39]: from sklearn.neighbors import KNeighborsClassifier  
classifier=KNeighborsClassifier(n_neighbors=11, metric="manhattan")  
classifier.fit(X_train_fs, y_train)
```

```
[39]: KNeighborsClassifier(metric='manhattan', n_neighbors=11)
```

```
[386]: y_pred=classifier.predict(X_test_fs)  
print(y_pred)
```

```
[1 1 1 0 0 0 0 0 0 1 1 1 1 0 1 1 1 1 0 1 1 1 1 0 0 0 1 1 1 1 1 0 0 0 1 0 0  
0 1 0 0 1 1 1 1 0 0 1 0 1 1 0 0 1 1 0 0 0 0 0 1 1 1 1 0 1 1 1 0 1 0 1 1 0 1  
1 1 1 1 1 1 1 1 0 0 1 1 1 1 0 1 0 0 0 1 0 1 0 0 0 0 1 1 1 1 0 1 0 0 0 1 0  
1 1 1 0 1 0 1 1 1 1 0 1 1 0 1 0 1 0 1 1 1 1 0 1 1 1 0 0 1 0 1 0 0 1 1  
0 0 0 1 0 0 1 0 1 1 1 0 1 0 1 0 1 0 1 0 1 1 0 0 1 1 1 0 0 1 1 1 0 1 1  
1 0 1 1 0 1 0 1 1 1 0 0 1 0 1 1 0 1 0 0 0 1 0 1 1 0 1 0 0 1 0 1 1 1 0  
0 1 1 1 0 1 1 1 0 1 1 1 0 1 1 0 1 0 1 0 0 0 0 1 0 1 0 1 0 0 0 1 0 0 1 0  
1 0 0 1 0 0 0 0 0 1 0 0 1 1 1 1]
```

```
[387]: from sklearn import metrics  
acc=metrics.accuracy_score(y_test, y_pred)  
print("accuracy: %.2f\n\n" % (acc))  
cm=metrics.confusion_matrix(y_test,y_pred)  
print("confusion Matrix: ")  
print(cm, "\n\n")  
print("-----")  
result=metrics.classification_report(y_test, y_pred)  
print("Classification Report:\n")  
print(result)
```

accuracy:0.87

confusion Matrix:

```
[[103  18]  
 [ 19 136]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.85	0.85	121
1	0.88	0.88	0.88	155
accuracy			0.87	276
macro avg	0.86	0.86	0.86	276

	weighted avg	0.87	0.87	0.87	276
--	--------------	------	------	------	-----

4.1 other k

```
[362]: for k in range(1, 12):
    classifier = KNeighborsClassifier(n_neighbors=k, metric="manhattan")
    classifier.fit(X_train_s, y_train)

    y_pred_train = classifier.predict(X_train_s)
    accuracy_train = metrics.accuracy_score(y_train, y_pred_train)

    y_pred=classifier.predict(X_test_fs)
    accuracy_test = metrics.accuracy_score(y_test, y_pred)

    print(f"n_neighbors = {k}, Train Accuracy = {accuracy_train:.2f}")
    print(f"n_neighbors = {k}, Test Accuracy = {accuracy_test:.2f}")
    print("")
```

n_neighbors = 1, Train Accuracy = 1.00
n_neighbors = 1, Test Accuracy = 0.84

n_neighbors = 2, Train Accuracy = 0.90
n_neighbors = 2, Test Accuracy = 0.79

n_neighbors = 3, Train Accuracy = 0.91
n_neighbors = 3, Test Accuracy = 0.85

n_neighbors = 4, Train Accuracy = 0.90
n_neighbors = 4, Test Accuracy = 0.83

n_neighbors = 5, Train Accuracy = 0.90
n_neighbors = 5, Test Accuracy = 0.84

n_neighbors = 6, Train Accuracy = 0.88
n_neighbors = 6, Test Accuracy = 0.85

n_neighbors = 7, Train Accuracy = 0.90
n_neighbors = 7, Test Accuracy = 0.84


n_neighbors = 8, Train Accuracy = 0.89
n_neighbors = 8, Test Accuracy = 0.83

n_neighbors = 9, Train Accuracy = 0.89
n_neighbors = 9, Test Accuracy = 0.84

n_neighbors = 10, Train Accuracy = 0.88
n_neighbors = 10, Test Accuracy = 0.86

```
n_neighbors = 11, Train Accuracy = 0.88  
n_neighbors = 11, Test Accuracy = 0.87
```

```
[407]: ax=sns.heatmap(cm, annot=True, fmt='d', cmap='flare')  
plt.title('Confusion Matrix')  
plt.xlabel('Predicted Class')  
plt.ylabel('True Class')  
plt.show()
```



5 Decision Tree

```
[40]: from sklearn.tree import DecisionTreeClassifier  
tree_classifier=DecisionTreeClassifier(criterion="gini", random_state=0)  
tree_classifier.fit(X_train_fs, y_train)
```

```
[40]: DecisionTreeClassifier(random_state=0)
```

```
[393]: tree_y_pred=tree_classifier.predict(X_test_fs)
print(tree_y_pred)

[1 1 1 0 1 0 1 0 0 0 1 1 1 1 0 1 1 1 1 0 1 1 0 0 0 0 1 1 1 1 1 1 0 0 0 1 0 0 0 1 1 1 1 1 1 0 0 0 1 1 1 1 1 1 0 1 1 1 1 0 1 0 1 1 1 1
```

```
[394]: from sklearn import metrics
tree_acc=metrics.accuracy_score(y_test, tree_y_pred)
print("accuracy: %.2f\n\n" % (tree_acc))
tree_cm=metrics.confusion_matrix(y_test,tree_y_pred)
print("confusion Matrix: ")
print(tree_cm, "\n\n")
print("-----")
tree_result=metrics.classification_report(y_test, tree_y_pred)
print("Classification Report:\n")
print(tree_result)
```

accuracy: 0.80

confusion Matrix:


```
[[ 98  23]
 [ 32 123]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.81	0.78	121
1	0.84	0.79	0.82	155
accuracy			0.80	276
macro avg	0.80	0.80	0.80	276
weighted avg	0.80	0.80	0.80	276

```
[395]: tree_ax=sns.heatmap(tree_cm, annot=True, fmt='d', cmap='flare')
plt.title('Confusion Matrix')
plt.xlabel('Predicted Class')
plt.ylabel('True Class')
```

```
plt.show()
```



5.1 sym

```
[41]: from sklearn import svm  
svm_classifier=svm.SVC(kernel="rbf")  
svm_classifier.fit(X_train_fs, y_train)
```

[41]: SVC()

```
[46]: svm_y_pred=svm_classifier.predict(X_test_fs)
      print(svm_y_pred)
```

```
0 1 1 1 0 1 1 0 1 1 0 1 1 0 1 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 0 0  
1 0 1 1 0 0 0 0 0 1 0 0 1 1 1 1 1]
```

```
[47]: from sklearn import metrics  
svm_acc=metrics.accuracy_score(y_test, svm_y_pred)  
print("accuracy: %.2f\n\n" % (svm_acc))  
svm_cm=metrics.confusion_matrix(y_test,svm_y_pred)  
print("confusion Matrix: ")  
print(svm_cm, "\n\n")  
print("-----")  
svm_result=metrics.classification_report(y_test, svm_y_pred)  
print("Classification Report:\n")  
print(svm_result)
```

accuracy:0.88


confusion Matrix:

```
[[102 19]  
 [ 13 142]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.84	0.86	121
1	0.88	0.92	0.90	155
accuracy			0.88	276
macro avg	0.88	0.88	0.88	276
weighted avg	0.88	0.88	0.88	276

```
[48]: svm_ax=sns.heatmap(svm_cm, annot=True, fmt='d', cmap='flare')  
plt.title('Confusion Matrix')  
plt.xlabel('Predicted Class')  
plt.ylabel('True Class')  
  
plt.show()
```



[]: