

jlzcqlna6

September 6, 2025

### 0.0.1 Step 1: Import Libraries

```
[11]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from scipy import stats

import warnings
warnings.filterwarnings("ignore")

sns.set(style="whitegrid")
```

```
[12]: # show all columns
pd.set_option('display.max_column',None)

# present scientific notation (optional)
pd.set_option('display.float_format','{:.2f}'.format)
```

### 0.0.2 Step 2: Load and Inspect Dataset

```
[13]: df = pd.read_csv(r"C:\Users\atulm\Desktop\Data Analytics\Python\Project\New_U
↳folder\retail_sales_50k.csv")
df.head()
```

```
[13]:   Order_ID Customer_ID  Gender  Age  Country Product_Category  Quantity \
0      10001      C00861    Male   40  Germany        Fashion         7
1      10002      C03773  Female   32    India        Home         7
2      10003      C03093    Male   28    India        Home         4
3      10004      C00467  Female   38     UAE        Home         4
4      10005      C04427    Male   56      UK  Groceries         7

      Unit_Price  Discount          Order_Date Payment_Method \
0        5232      0.00  2023-01-01 00:00:00.000000000  Credit Card
1        8563      0.00  2023-01-01 00:21:01.465229304  Debit Card
2        3369      0.10  2023-01-01 00:42:02.930458609  Debit Card
```

```
3      6796      0.10  2023-01-01 01:03:04.395687913      Debit Card
4      2949      0.15  2023-01-01 01:24:05.860917218      COD
```

	Ad_Campaign	Returned	Total_Sales
0	B	No	36624.00
1	A	No	59941.00
2	A	No	12128.40
3	A	No	24465.60
4	A	No	17546.55

```
[14]: # Quick look
print("\nshape:", df.shape)
print("\ninfo:")
df.info()
```

```
shape: (50000, 14)

info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Order_ID         50000 non-null   int64  
 1   Customer_ID     50000 non-null   object  
 2   Gender           50000 non-null   object  
 3   Age              50000 non-null   int64  
 4   Country          50000 non-null   object  
 5   Product_Category 50000 non-null   object  
 6   Quantity         50000 non-null   int64  
 7   Unit_Price       50000 non-null   int64  
 8   Discount          50000 non-null   float64 
 9   Order_Date        50000 non-null   object  
 10  Payment_Method   50000 non-null   object  
 11  Ad_Campaign      50000 non-null   object  
 12  Returned          50000 non-null   object  
 13  Total_Sales      50000 non-null   float64 
dtypes: float64(2), int64(4), object(8)
memory usage: 5.3+ MB
```

```
[15]: # Summary Statistics
print('summary:')
df.describe()
```

```
summary:
```

```
[15]:    Order_ID      Age  Quantity  Unit_Price  Discount  Total_Sales
count  50000.00  50000.00  50000.00  50000.00  50000.00  50000.00
mean   35000.50   38.50     5.00    5058.98    0.06   23590.88
std    14433.90   12.12     2.59    2857.42    0.07   19358.09
min    10001.00   18.00     1.00    100.00    0.00     89.60
25%   22500.75   28.00     3.00    2580.00    0.00   7671.90
50%   35000.50   39.00     5.00    5068.00    0.05   18160.10
75%   47500.25   49.00     7.00    7549.00    0.10   35309.10
max   60000.00   59.00     9.00   9999.00    0.20  89991.00
```

### 0.0.3 Step 3: Data Cleaning & Preprocessing

```
[16]: df["Order_Date"]
```

```
[16]: 0      2023-01-01 00:00:00.000000000
1      2023-01-01 00:21:01.465229304
2      2023-01-01 00:42:02.930458609
3      2023-01-01 01:03:04.395687913
4      2023-01-01 01:24:05.860917218
...
49995  2024-12-30 22:35:54.139082784
49996  2024-12-30 22:56:55.604312088
49997  2024-12-30 23:17:57.069541392
49998  2024-12-30 23:38:58.534770696
49999  2024-12-31 00:00:00.000000000
Name: Order_Date, Length: 50000, dtype: object
```

```
[17]: # Change Format remove text unes..
df["Order_Date"] = pd.to_datetime(df["Order_Date"])
df["Order_Date"] = df["Order_Date"].dt.strftime("%Y-%m-%d")
df["Order_Date"]
```

```
[17]: 0      2023-01-01
1      2023-01-01
2      2023-01-01
3      2023-01-01
4      2023-01-01
...
49995  2024-12-30
49996  2024-12-30
49997  2024-12-30
49998  2024-12-30
49999  2024-12-31
Name: Order_Date, Length: 50000, dtype: object
```

```
[18]: # Check Missing Values
print("\nMissing Values :\n")
```

```
df.isnull().sum()
```

Missing Values :

```
[18]: Order_ID      0  
Customer_ID      0  
Gender           0  
Age              0  
Country          0  
Product_Category 0  
Quantity          0  
Unit_Price       0  
Discount          0  
Order_Date       0  
Payment_Method   0  
Ad_Campaign      0  
Returned          0  
Total_Sales      0  
dtype: int64
```

```
[19]: df.head()
```

```
[19]:    Order_ID Customer_ID  Gender  Age  Country Product_Category  Quantity  \  
0      10001      C00861    Male    40  Germany        Fashion         7  
1      10002      C03773  Female   32   India        Home          7  
2      10003      C03093    Male    28   India        Home          4  
3      10004      C00467  Female   38    UAE        Home          4  
4      10005      C04427    Male    56     UK        Groceries        7  
  
      Unit_Price  Discount  Order_Date Payment_Method Ad_Campaign Returned  \  
0        5232     0.00  2023-01-01    Credit Card        B      No  
1        8563     0.00  2023-01-01    Debit Card        A      No  
2        3369     0.10  2023-01-01    Debit Card        A      No  
3        6796     0.10  2023-01-01    Debit Card        A      No  
4        2949     0.15  2023-01-01        COD        A      No  
  
      Total_Sales  
0      36624.00  
1      59941.00  
2      12128.40  
3      24465.60  
4      17546.55
```

#### 0.0.4 Step 4: Univariate Analysis

```
[21]: columns = ['Year', 'Age', 'Gender', 'Country',
                 'Product_Category', 'Quantity', 'Payment_Method', 'Ad_Campaign']

for col in columns:
    if col in df.columns:
        print(f"\nUnique Count of {col} Column -: {df[col].nunique()}")
        print("\nValue Counts of the Column -:\n", df[col].value_counts())
        print("\n-----")
    else:
        print(f"\nColumn '{col}' not found in the DataFrame!")
        print("\n-----")
```

Column 'Year' not found in the DataFrame!

-----

Unique Count of Age Column -: 42

Value Counts of the Column -:

Age	Count
40	1284
18	1247
50	1241
28	1239
44	1236
54	1232
37	1219
32	1219
48	1217
56	1215
58	1214
29	1213
31	1211
33	1210
19	1207
39	1202
30	1201
55	1198
43	1189
47	1187
25	1185
34	1184
22	1180
57	1179
24	1173

```
51    1171
49    1170
26    1169
38    1168
46    1168
21    1167
20    1166
52    1166
35    1165
59    1163
27    1161
42    1157
53    1155
41    1152
23    1144
36    1143
45    1133
Name: count, dtype: int64
```

---

Unique Count of Gender Column :- 2

Value Counts of the Column :-  
Gender  
Female 25114  
Male 24886  
Name: count, dtype: int64

---

Unique Count of Country Column :- 5

Value Counts of the Column :-  
Country  
India 19877  
UAE 10053  
USA 9933  
UK 5119  
Germany 5018  
Name: count, dtype: int64

---

Unique Count of Product\_Category Column :- 5

Value Counts of the Column :-  
Product\_Category

```
Sports      10111
Home       10000
Electronics 9973
Fashion     9969
Groceries   9947
Name: count, dtype: int64
```

---

Unique Count of Quantity Column -: 9

Value Counts of the Column -:

Quantity	Count
3	5664
2	5635
8	5628
9	5598
6	5525
5	5513
1	5491
4	5483
7	5463

```
Name: count, dtype: int64
```

---

Unique Count of Payment\_Method Column -: 4

Value Counts of the Column -:

Payment_Method	Count
UPI	12629
Debit Card	12469
COD	12469
Credit Card	12433

```
Name: count, dtype: int64
```

---

Unique Count of Ad\_Campaign Column -: 2

Value Counts of the Column -:

Ad_Campaign	Count
A	25093
B	24907

```
Name: count, dtype: int64
```

---

```
[22]: df.head(1)
```

```
[22]:   Order_ID Customer_ID Gender  Age  Country Product_Category  Quantity \
0      10001      C00861    Male   40  Germany          Fashion       7
                                               Unit_Price  Discount  Order_Date Payment_Method Ad_Campaign Returned \
0        5232       0.00  2023-01-01     Credit Card           B        No
                                               Total_Sales
0      36624.00
```

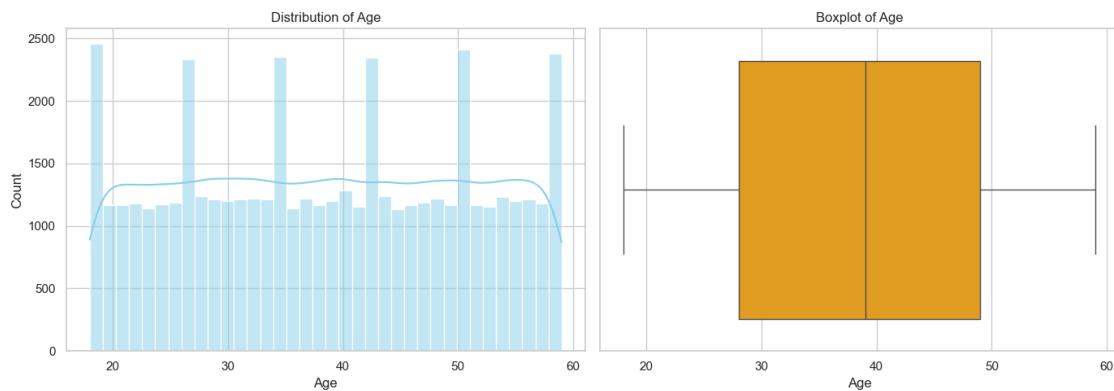
```
[23]: numerical_cols = ['Age', 'Quantity', 'Unit_Price', 'Discount', 'Total_Sales']
```

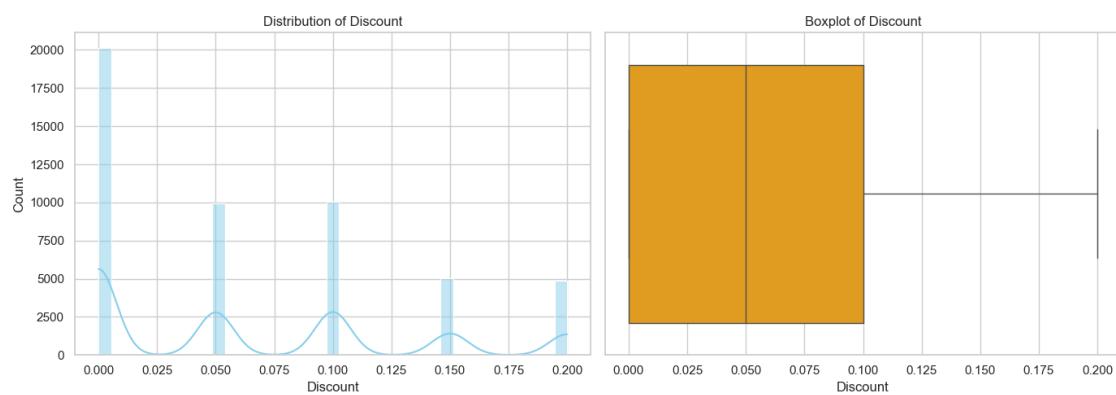
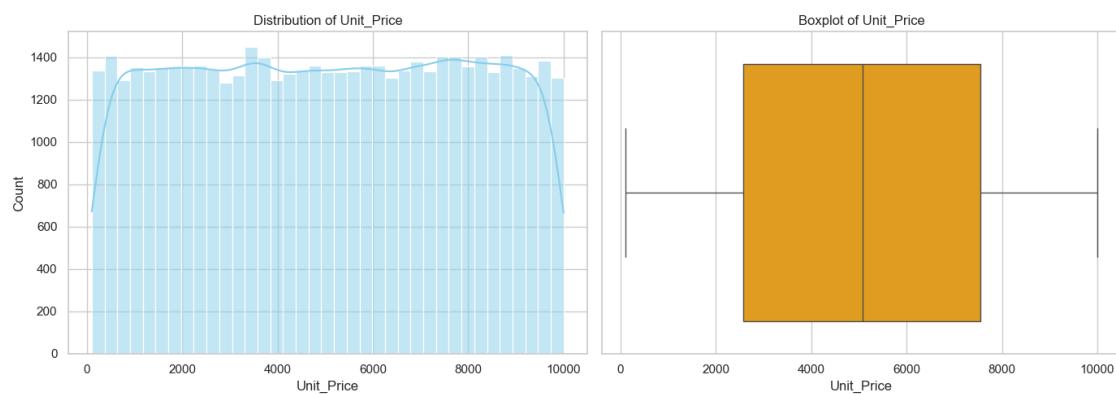
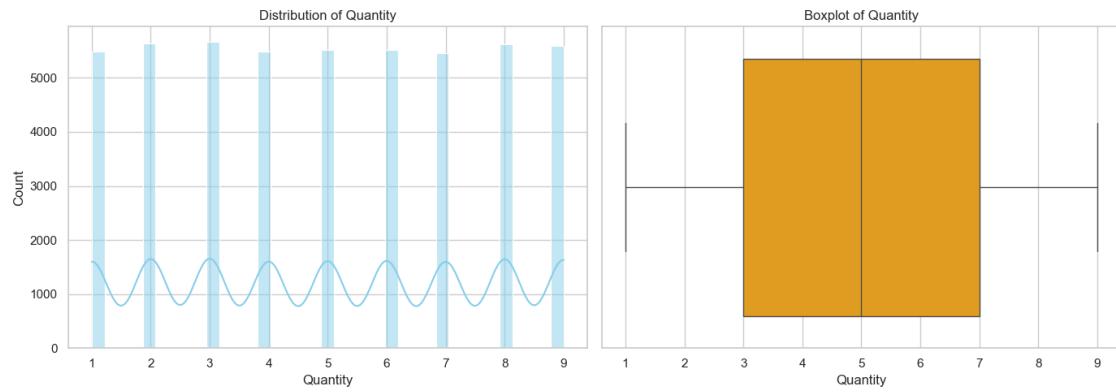
```
for col in numerical_cols:
    plt.figure(figsize=(14, 5))

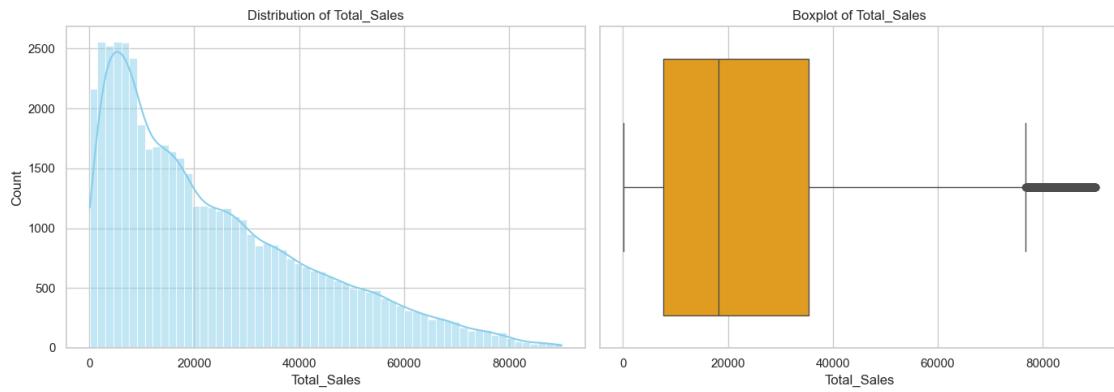
    # Histogram
    plt.subplot(1, 2, 1)
    sns.histplot(df[col], kde=True, color='skyblue')
    plt.title(f'Distribution of {col}')

    # Boxplot
    plt.subplot(1, 2, 2)
    sns.boxplot(x=df[col], color='orange')
    plt.title(f'Boxplot of {col}')

plt.tight_layout()
plt.show()
```

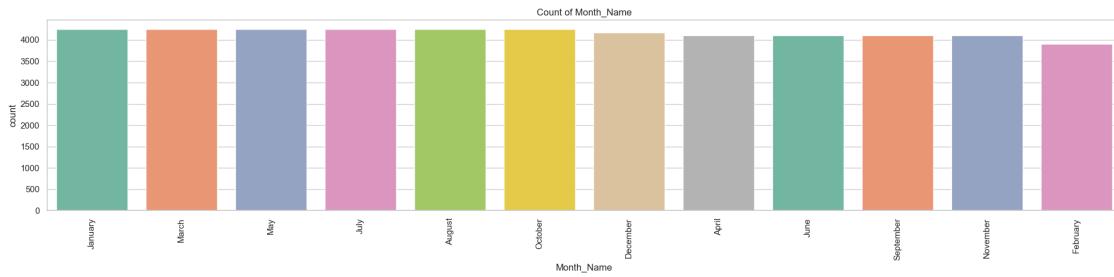


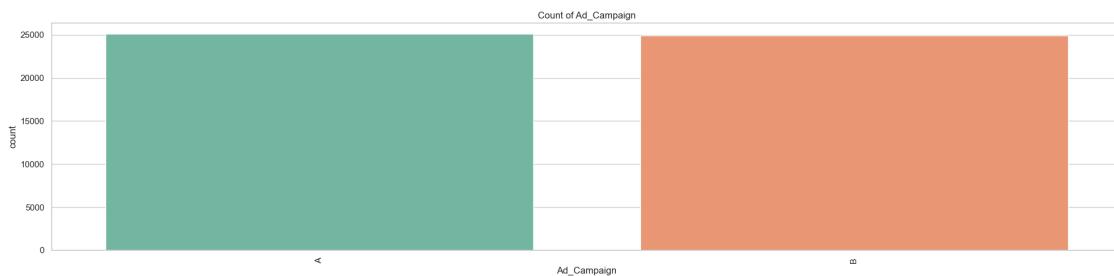
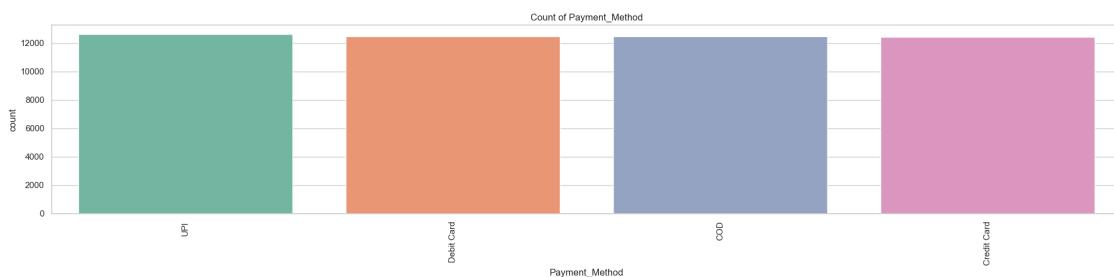
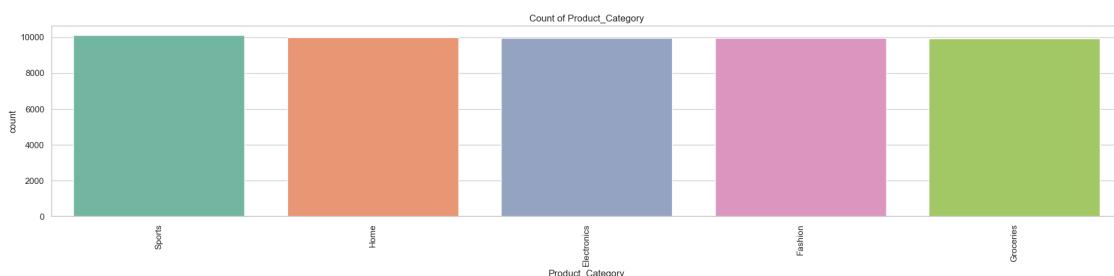
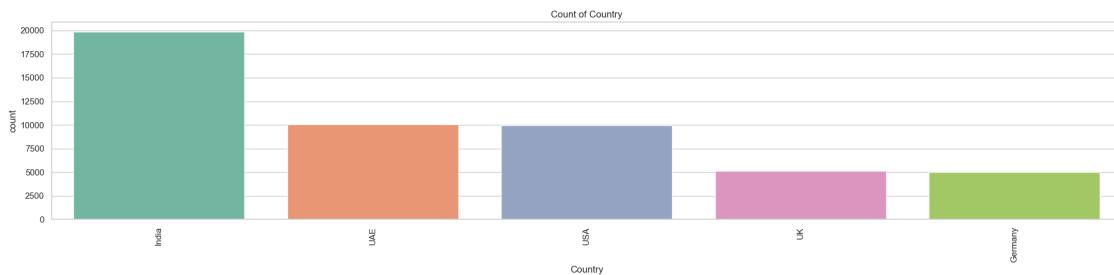




```
[ ]: categorical_cols = ['Month_Name', 'Gender', 'Country',
                       'Product_Category', 'Payment_Method',
                       'Ad_Campaign', 'Returned']

for col in categorical_cols:
    plt.figure(figsize=(20, 5))
    sns.countplot(x=col, data=df, order=df[col].value_counts().index, palette='Set2')
    plt.xticks(rotation=90)
    plt.title(f'Count of {col}')
    plt.tight_layout()
    plt.show()
```







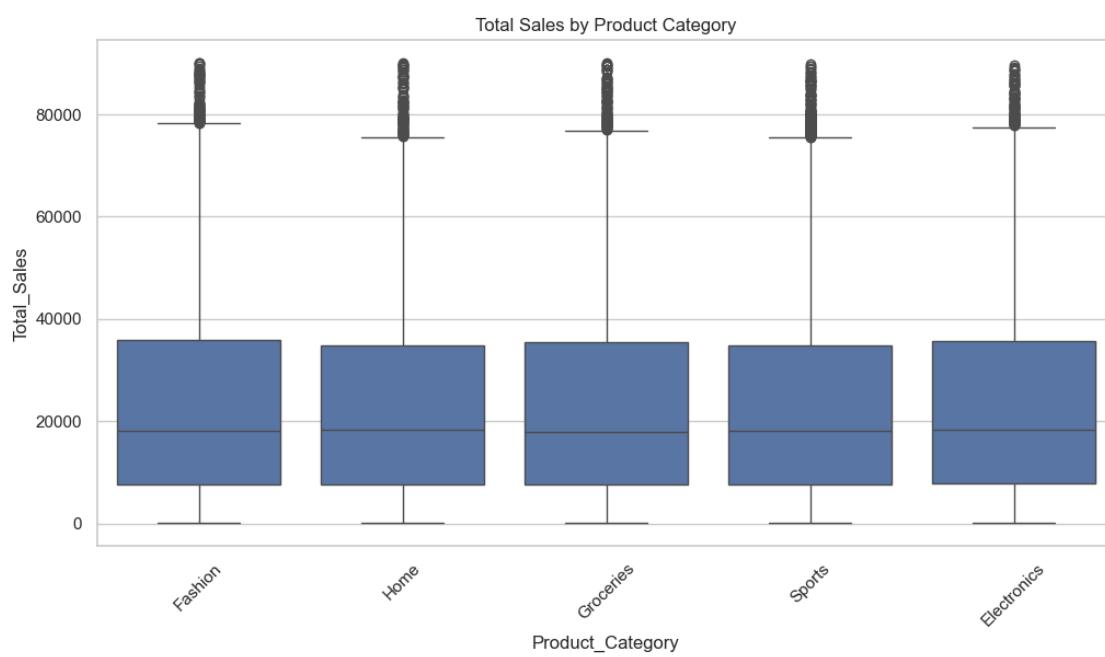
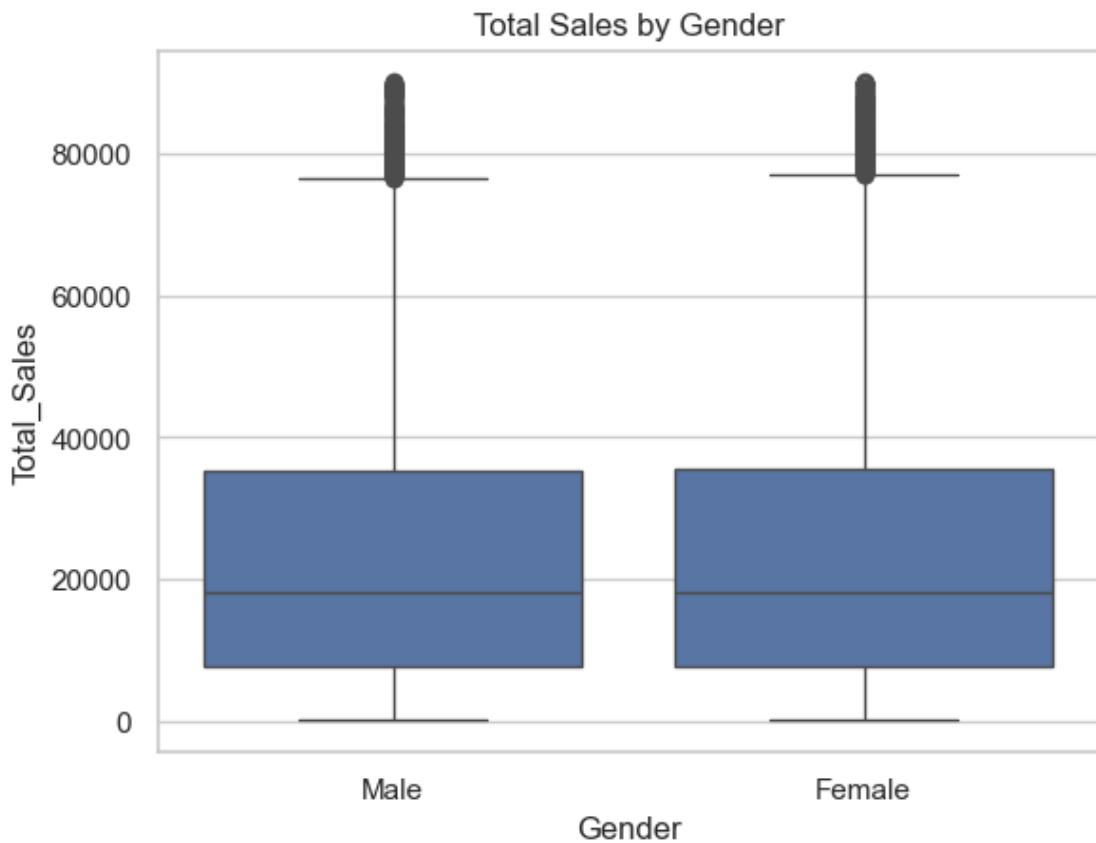
### 0.0.5 Step 5: Bivariate Analysis

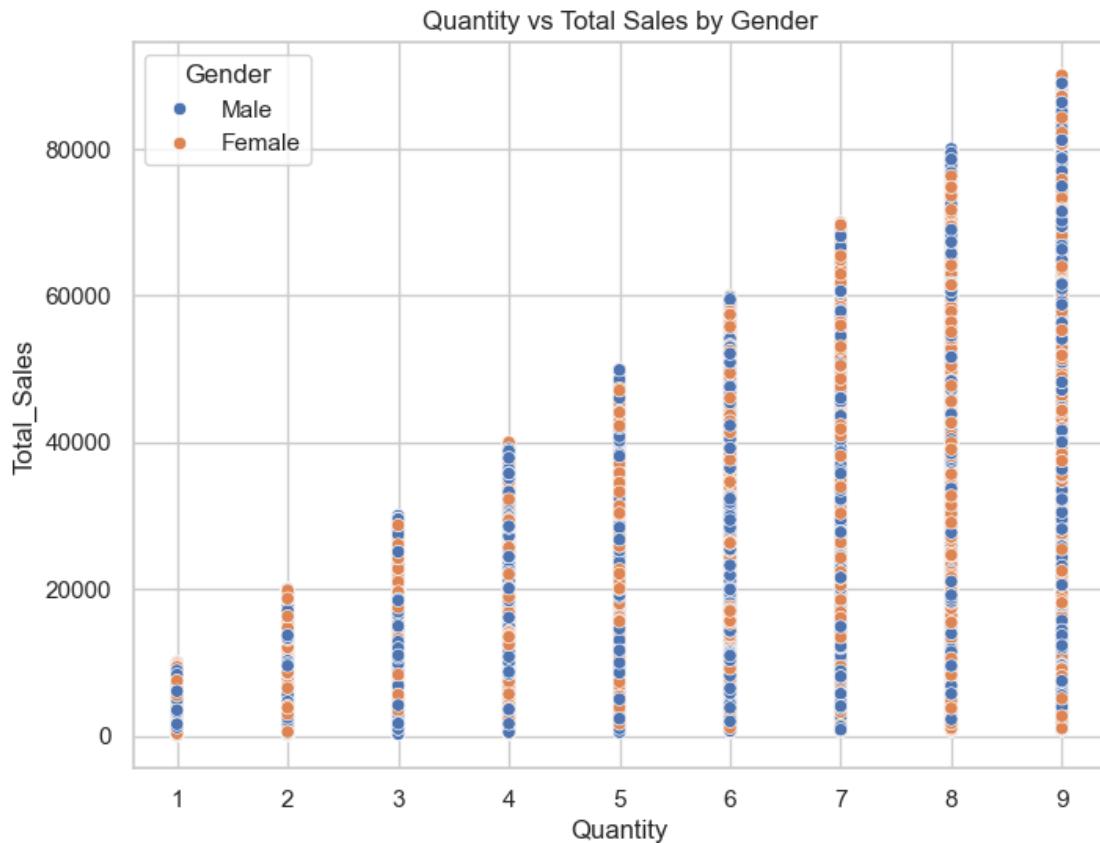
Definition: Bivariate Analysis ka matlab hai do variables ke beech relationship ko samajhnna.

```
[ ]: # Revenue by Gender
sns.boxplot(x='Gender', y='Total_Sales', data=df)
plt.title('Total Sales by Gender')
plt.show()

# Revenue by Product Category
plt.figure(figsize=(12,6))
sns.boxplot(x='Product_Category', y='Total_Sales', data=df)
plt.xticks(rotation=45)
plt.title('Total Sales by Product Category')
plt.show()

# Quantity vs Revenue (Scatter)
plt.figure(figsize=(8,6))
sns.scatterplot(x='Quantity', y='Total_Sales', data=df, hue='Gender')
plt.title("Quantity vs Total Sales by Gender")
plt.show()
```





- Top & bottom products by total revenue
- Top & bottom products by total quantity sold
- Top & bottom products by total profit (assuming Profit = Revenue - Cost)

## 0.1 Top & Bottom Products Analysis – Full Code

```
[ ]: df.head(1)
```

```
[ ]: Order_ID Customer_ID Gender Age Country Product_Category Quantity \
0 10001 C00861 Male 40 Germany Fashion 7
Unit_Price Discount Order_Date Payment_Method Ad_Campaign Returned \
0 5232 0.00 2023-01-01 Credit Card B No
Total_Sales Year Month_Name
0 36624.00 2023 January
```

```
[ ]: # Step 1: Create Unit_Cost and Profit
df['Unit_Cost'] = df['Unit_Price'] * 0.7      # assume 70% cost
df['Profit'] = (df['Unit_Price'] - df['Unit_Cost']) * df['Quantity']
```

```

# Step 2: Group by Product_Category
product_perf = df.groupby('Product_Category').agg({
    'Quantity': 'sum',
    'Total_Sales': 'sum',
    'Profit': 'sum'
}).sort_values(by='Profit', ascending=False)

# Step 3: Show the full performance table
print("\n Product Performance Summary:")
print(product_perf)

```

Product Performance Summary:

Product_Category	Quantity	Total_Sales	Profit
Electronics	50168	238001936.80	76289367.30
Sports	50242	237584019.20	76223923.20
Fashion	49945	235476486.00	75491202.90
Home	49915	234025872.70	75115932.30
Groceries	49777	234455479.45	75091210.80

```
[ ]: # Top 5 Products by Total Sales
top5_sales = product_perf.sort_values(by='Total_Sales', ascending=False).head(5)

print(" Top 5 Products by Total Sales:")
print(top5_sales)
```

Top 5 Products by Total Sales:

Product_Category	Quantity	Total_Sales	Profit
Electronics	50168	238001936.80	76289367.30
Sports	50242	237584019.20	76223923.20
Fashion	49945	235476486.00	75491202.90
Groceries	49777	234455479.45	75091210.80
Home	49915	234025872.70	75115932.30

```
[ ]: # Bottom 5 Products by Total Sales (Revenue)
bottom_sales = product_perf.sort_values(by='Total_Sales', ascending=True) .
    ↵head(5)

print("\n Bottom 5 Products by Total Sales:")
print(bottom_sales)
```

Bottom 5 Products by Total Sales:

Product_Category	Quantity	Total_Sales	Profit
------------------	----------	-------------	--------

Home	49915	234025872.70	75115932.30
Groceries	49777	234455479.45	75091210.80
Fashion	49945	235476486.00	75491202.90
Sports	50242	237584019.20	76223923.20
Electronics	50168	238001936.80	76289367.30

```
[ ]: # Top 5 Products by Quantity Sold
top_quantity = product_perf.sort_values(by='Quantity', ascending=False).head(5)
print("\n Top 5 Products by Quantity Sold:")
print(top_quantity)
```

Top 5 Products by Quantity Sold:

	Quantity	Total_Sales	Profit
Product_Category			
Sports	50242	237584019.20	76223923.20
Electronics	50168	238001936.80	76289367.30
Fashion	49945	235476486.00	75491202.90
Home	49915	234025872.70	75115932.30
Groceries	49777	234455479.45	75091210.80

```
[ ]: # Bottom 5 Products by Quantity Sold
bottom_quantity = product_perf.sort_values(by='Quantity', ascending=True).
    ↪head(5)
print("\n Bottom 5 Products by Quantity Sold:")
print(bottom_quantity)
```

Bottom 5 Products by Quantity Sold:

	Quantity	Total_Sales	Profit
Product_Category			
Groceries	49777	234455479.45	75091210.80
Home	49915	234025872.70	75115932.30
Fashion	49945	235476486.00	75491202.90
Electronics	50168	238001936.80	76289367.30
Sports	50242	237584019.20	76223923.20

```
[ ]: # Top 5 Products by Profit
top_profit = product_perf.sort_values(by='Profit', ascending=False).head(5)
print("\n Top 5 Products by Profit:")
print(top_profit)
```

Top 5 Products by Profit:

	Quantity	Total_Sales	Profit
Product_Category			
Electronics	50168	238001936.80	76289367.30
Sports	50242	237584019.20	76223923.20

```

Fashion          49945 235476486.00 75491202.90
Home            49915 234025872.70 75115932.30
Groceries       49777 234455479.45 75091210.80

```

```
[ ]: # Bottom 5 Products by Profit
bottom_profit = product_perf.sort_values(by='Profit', ascending=True).head(5)
print("\n Bottom 5 Products by Profit:")
print(bottom_profit)
```

Bottom 5 Products by Profit:

	Quantity	Total_Sales	Profit
Product_Category			
Groceries	49777	234455479.45	75091210.80
Home	49915	234025872.70	75115932.30
Fashion	49945	235476486.00	75491202.90
Sports	50242	237584019.20	76223923.20
Electronics	50168	238001936.80	76289367.30

```
[ ]: # Category-wise Revenue/Profit Share
category_share = product_perf[['Total_Sales', 'Profit']].apply(lambda x: x/x.sum() * 100)
print(category_share)
```

	Total_Sales	Profit
Product_Category		
Electronics	20.18	20.17
Sports	20.14	20.15
Fashion	19.96	19.96
Home	19.84	19.86
Groceries	19.88	19.85

```
[ ]: # Revenue vs Quantity Correlation
product_perf[['Quantity', 'Total_Sales', 'Profit']].corr()
```

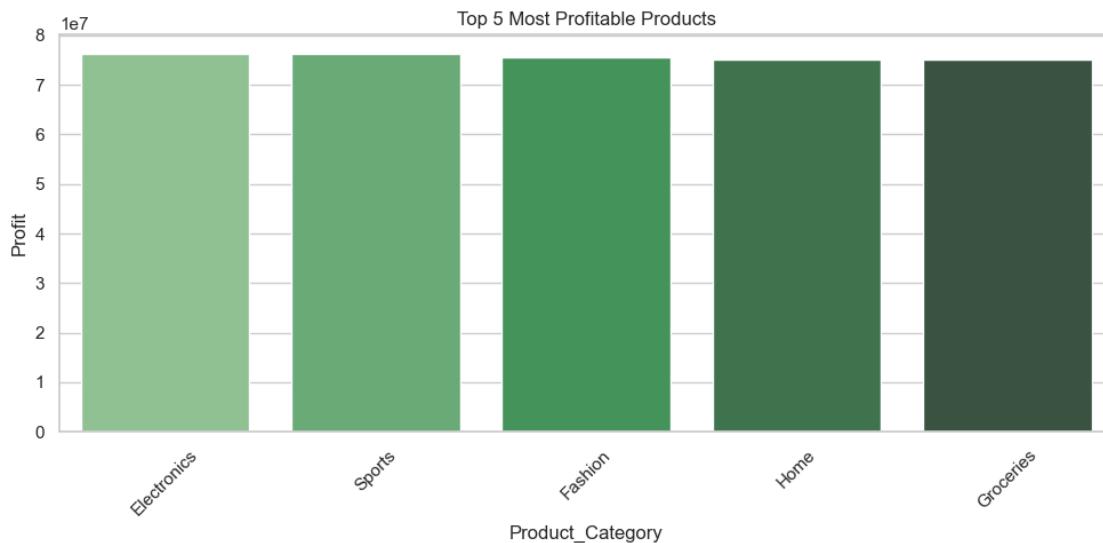
	Quantity	Total_Sales	Profit
Quantity	1.00	0.92	0.95
Total_Sales	0.92	1.00	0.99
Profit	0.95	0.99	1.00

```
[ ]: gender_perf = df.groupby('Gender').agg({
    'Total_Sales':'sum',
    'Profit':'sum'
})
print(gender_perf)
```

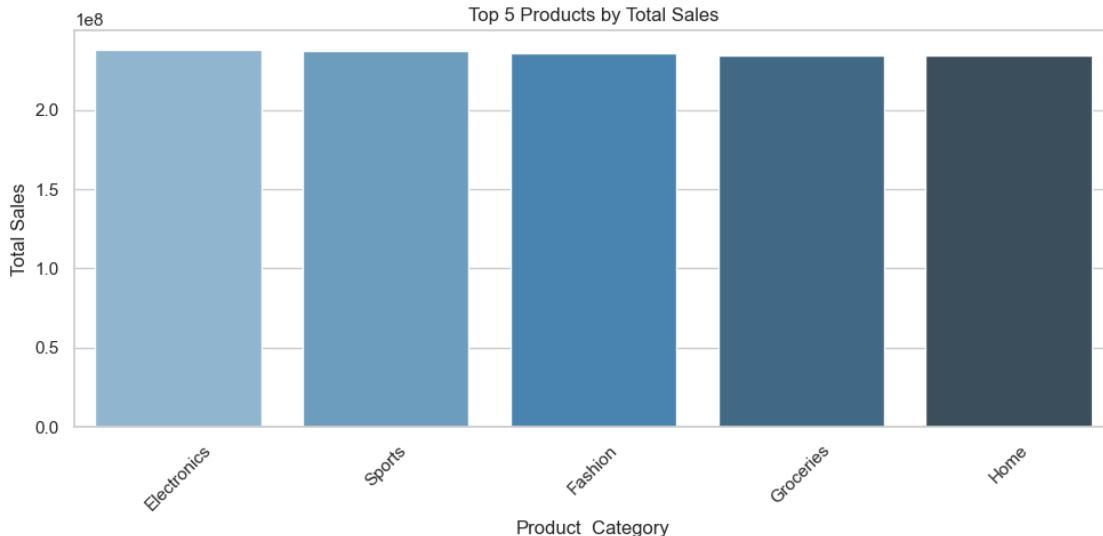
	Total_Sales	Profit
Gender		

```
Female 594913301.40 190727994.90
Male    584630492.75 187483641.60
```

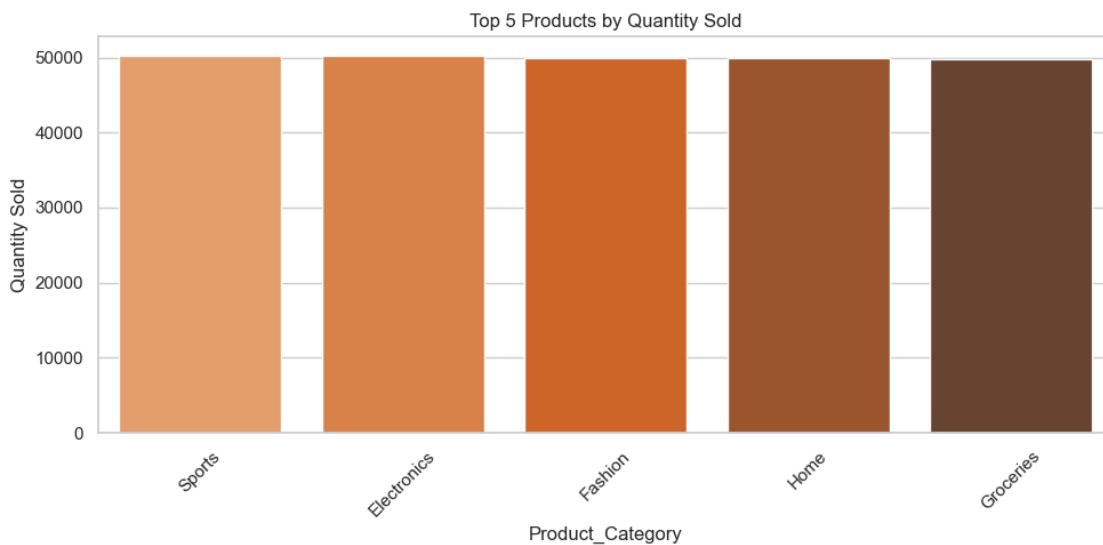
```
[ ]: # Top 5 Profitable Products
top_profit = product_perf.sort_values(by='Profit', ascending=False).head(5)
plt.figure(figsize=(10,5))
sns.barplot(x=top_profit.index, y=top_profit['Profit'], palette='Greens_d')
plt.title("Top 5 Most Profitable Products")
plt.ylabel("Profit")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[ ]: # Top 5 Products by Total Sales
top_sales = product_perf.sort_values(by='Total_Sales', ascending=False).head(5)
plt.figure(figsize=(10,5))
sns.barplot(x=top_sales.index, y=top_sales['Total_Sales'], palette='Blues_d')
plt.title("Top 5 Products by Total Sales")
plt.ylabel("Total Sales")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[ ]: # Top 5 Products by Quantity Sold
top_quantity = product_perf.sort_values(by='Quantity', ascending=False).head(5)
plt.figure(figsize=(10,5))
sns.barplot(x=top_quantity.index, y=top_quantity['Quantity'], palette='Oranges_d')
plt.title("Top 5 Products by Quantity Sold")
plt.ylabel("Quantity Sold")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[ ]: df['Order_Date'] = pd.to_datetime(df['Order_Date'], errors='coerce')
```

```
[ ]: # Monthly Total Sales Trend
df['Month_Year'] = df['Order_Date'].dt.to_period('M')
monthly_sales = df.groupby('Month_Year')['Total_Sales'].sum()

plt.figure(figsize=(12,6))
monthly_sales.plot(marker='o')
plt.title("Monthly Total Sales Trend")
plt.ylabel("Total Sales")
plt.xlabel("Month")
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[ ]: gender_perf[['Total_Sales', 'Profit']].plot(kind='bar', figsize=(8,5))
plt.title("Revenue & Profit by Gender")
plt.ylabel("Amount")
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```

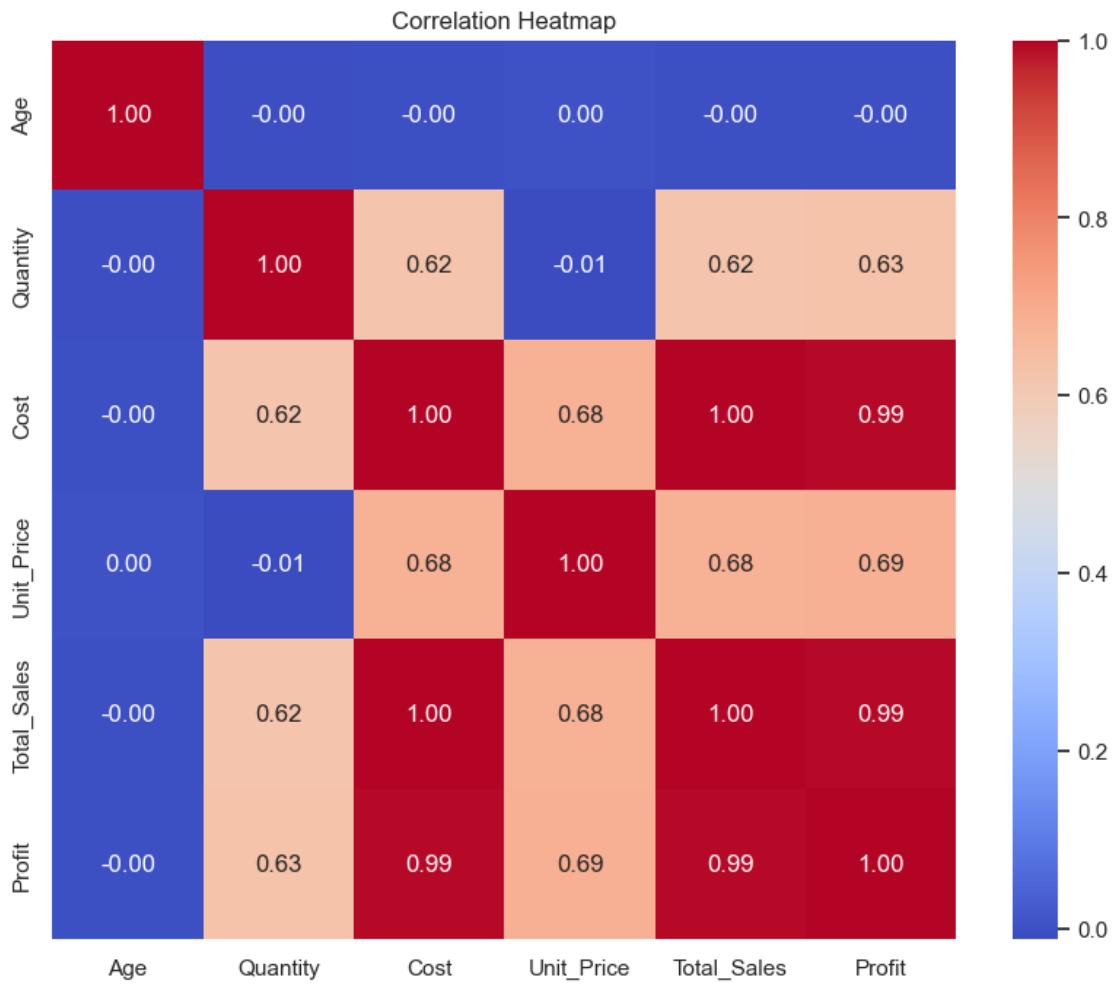


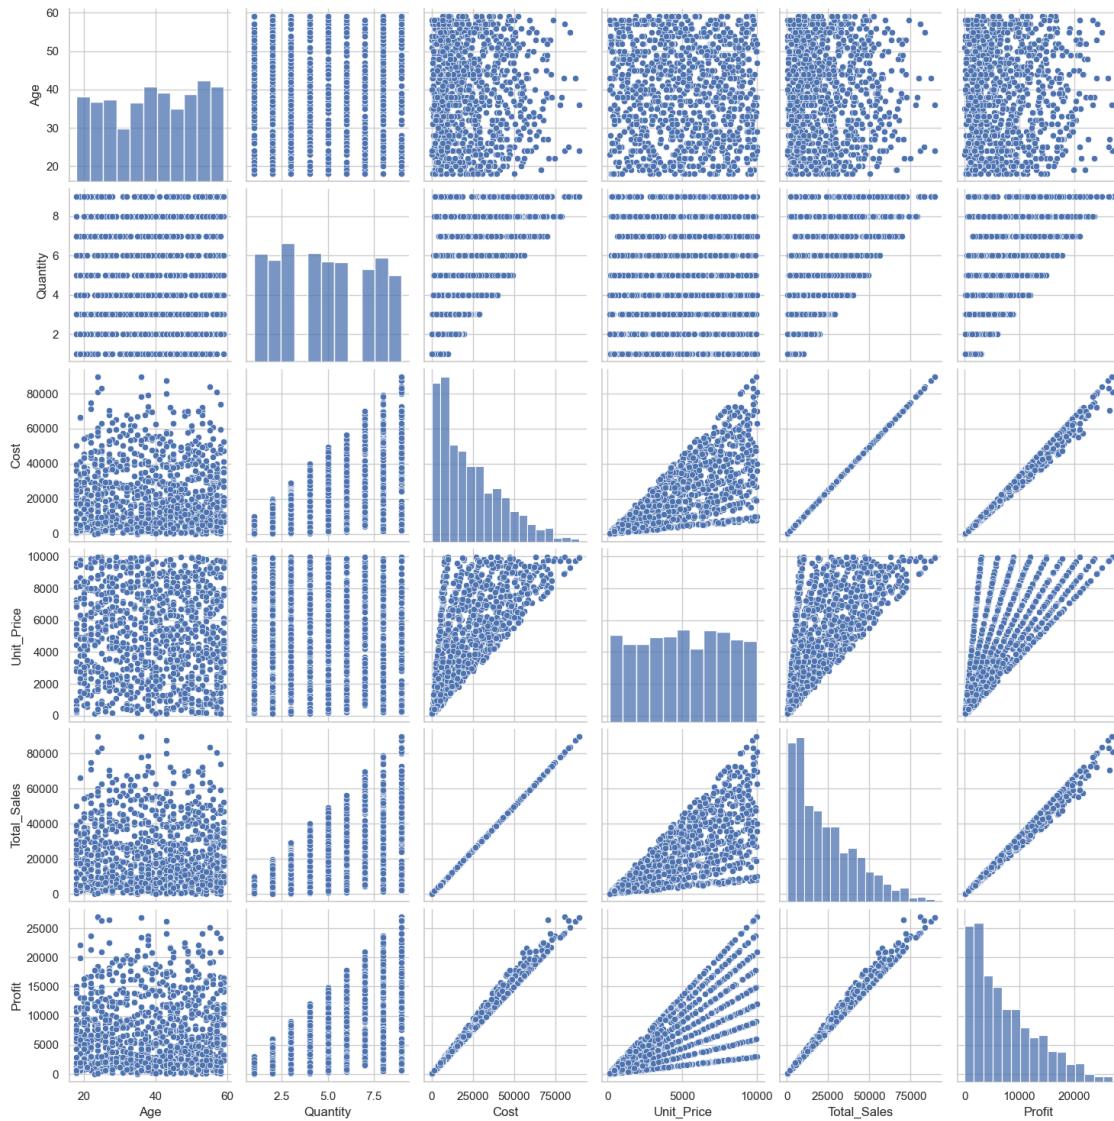
### 0.1.1 Step 6: Multivariate Analysis

```
[ ]: # Numerical columns
numerical_cols = ['Age', 'Quantity', 'Cost', 'Unit_Price', 'Total_Sales', ↴
                  'Profit']

# 1 Correlation Heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(df[numerical_cols].corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()

# 2 Pairplot (sample for large dataset)
sample_df = df[numerical_cols].sample(1000, random_state=1) # 1000 random rows
sns.pairplot(sample_df)
plt.show()
```





### 0.1.2 Step 7: Hypothesis Testing

```
[ ]: df.groupby('Gender')[['Total_Sales']].agg(['sum','mean','max','min'])
```

```
[ ]:      sum      mean      max      min
Gender
Female  594913301.40  23688.51  89991.00  97.85
Male    584630492.75  23492.34  89991.00  89.60
```

```
[ ]: # Male and Female Total_Sales
revenue_by_gender = df.groupby('Gender')[['Total_Sales']].sum().reset_index()
revenue_by_gender
```

```
[ ]: Gender Total_Sales
0 Female 594913301.40
1 Male 584630492.75

[ ]: # Grouping by Gender and Product_Category
subcat_by_gender = df.groupby(['Gender', 'Product_Category'])['Total_Sales'].
    ↪sum().reset_index()
subcat_by_gender = subcat_by_gender.sort_values(by='Total_Sales', ↪
    ↪ascending=False)

# Ranking per Gender
subcat_by_gender["Rnk"] = subcat_by_gender.groupby('Gender')["Total_Sales"].
    ↪rank(method='dense', ascending=False)

# Top 5 per Gender
top5_subcat_by_gender = subcat_by_gender[subcat_by_gender["Rnk"] <= 5].
    ↪sort_values(by='Gender')
top5_subcat_by_gender
```

```
[ ]: Gender Product_Category Total_Sales Rnk
0 Female Electronics 120577510.10 1.00
4 Female Sports 120082832.05 2.00
2 Female Groceries 118984478.40 3.00
1 Female Fashion 117639557.85 4.00
3 Female Home 117628923.00 5.00
6 Male Fashion 117836928.15 1.00
9 Male Sports 117501187.15 2.00
5 Male Electronics 117424426.70 3.00
8 Male Home 116396949.70 4.00
7 Male Groceries 115471001.05 5.00
```

```
[ ]: # Grouping by Gender and Product_Category
subcat_by_gender_qty = df.groupby(['Gender', 'Product_Category'])['Quantity'].
    ↪sum().reset_index()
subcat_by_gender_qty = subcat_by_gender_qty.sort_values(by='Quantity', ↪
    ↪ascending=False)

# Ranking per Gender
subcat_by_gender_qty["Rnk"] = subcat_by_gender_qty.
    ↪groupby('Gender')[["Quantity"]].rank(method='dense', ascending=False)

# Top 5 per Gender
top5_subcat_by_gender_qty = subcat_by_gender_qty[subcat_by_gender_qty["Rnk"] <= 5].
    ↪sort_values(by='Gender')
top5_subcat_by_gender_qty
```

```
[ ]: Gender Product_Category Quantity Rnk
0 Female Electronics 25297 1.00
4 Female Sports 25279 2.00
2 Female Groceries 25141 3.00
1 Female Fashion 24899 4.00
3 Female Home 24854 5.00
8 Male Home 25061 1.00
6 Male Fashion 25046 2.00
9 Male Sports 24963 3.00
5 Male Electronics 24871 4.00
7 Male Groceries 24636 5.00
```

```
[ ]: df['Year'].unique()
```

```
[ ]: array([2023, 2024])
```

Z-Test : Compare Revenue Of the Year 2023 and 2024

Null Hypothesis (H0): The revenue of the year 2023 is equal to the revenue year 2024.

Alternative Hypothesis (H1) : The revenue 2023 is different from 2024.

```
[ ]: from scipy import stats

# Select revenue for each year
revenue_2023 = df[df['Year'] == 2023]['Total_Sales']
revenue_2024 = df[df['Year'] == 2024]['Total_Sales']

# Perform independent two-sample t-test
t_stat, p_value = stats.ttest_ind(revenue_2023, revenue_2024, equal_var=False)
# Welch's t-test

print(f"T-Test Statistic: {t_stat:.4f}")
print(f"P-Value: {p_value:.4f}")

# Interpretation
if p_value < 0.05:
    print("Reject Null Hypothesis: Revenue in 2023 and 2024 is significantly different.")
else:
    print("Accept Null Hypothesis: No significant difference in revenue between 2023 and 2024.")
```

T-Test Statistic: 0.0904

P-Value: 0.9280

Accept Null Hypothesis: No significant difference in revenue between 2023 and 2024.

```
[ ]: df.groupby(by = "Year")["Total_Sales"].agg(['sum', 'mean', 'max', 'min'])
```

```
[ ]:          sum      mean      max      min
Year
2023  589967593.95  23598.70  89991.00  96.30
2024  589576200.20  23583.05  89991.00  89.60
```

- Top-selling product per Year

```
[ ]: # Grouping by Year and Product_Category
Category_by_year = df.groupby(['Year', 'Product_Category'])['Quantity'].sum().
    ↪reset_index()

# Rank categories per year
Category_by_year["Rnk"] = Category_by_year.groupby('Year')["Quantity"].
    ↪rank(method='dense', ascending=False)

# Filter top 5 per year
top5_categories_per_year = Category_by_year[Category_by_year["Rnk"] <= 5].
    ↪sort_values(by='Year')

# Show final table
top5_categories_per_year
```

```
[ ]:   Year Product_Category  Quantity  Rnk
0  2023     Electronics    24851  4.00
1  2023        Fashion    25150  2.00
2  2023     Groceries    25186  1.00
3  2023        Home    24885  3.00
4  2023        Sports    24839  5.00
5  2024     Electronics    25317  2.00
6  2024        Fashion    24795  4.00
7  2024     Groceries    24591  5.00
8  2024        Home    25030  3.00
9  2024        Sports    25403  1.00
```

```
[ ]: from scipy import stats

# ANOVA: Revenue by Product Category
anova_data = [group['Total_Sales'] for name, group in df.
    ↪groupby('Product_Category')]
f_stat, p_val = stats.f_oneway(*anova_data)

print("ANOVA: Revenue by Product Category")
print(f"F-statistic = {f_stat:.4f}, P-value = {p_val:.4f}")

# Interpretation
```

```

if p_val < 0.05:
    print("Reject Null Hypothesis: Revenue significantly changed by Product\u202aCategory")
else:
    print("Accept Null Hypothesis: Revenue did not significantly change by\u202aProduct Category")

```

ANOVA: Revenue by Product Category  
F-statistic = 0.8026, P-value = 0.5233  
Accept Null Hypothesis: Revenue did not significantly change by Product Category

- Chi-Square Test: Relationship Between Customer Gender and Product Category

Null Hypothesis (H0) : There is no relationship between customer gender and product category.  
Alternative Hypothesis (H1) : There is a significant relationship between customer gender and product category.

```

[ ]: # Contingency table
contingency = pd.crosstab(df['Gender'], df['Product_Category'])
print("Contingency Table:\n", contingency)

# Chi-square test
chi2, p, dof, expected = chi2_contingency(contingency)

print(f"\nChi-Square Statistic: {chi2:.4f}")
print(f"Degrees of Freedom: {dof}")
print(f"P-Value: {p:.4f}")

# Interpretation
if p < 0.05:
    print(" Reject Null Hypothesis: Gender and Product Category are dependent.\n")
else:
    print(" Accept Null Hypothesis: No significant relationship between Gender\u202aand Product Category.")

```

Contingency Table:

	Product_Category	Electronics	Fashion	Groceries	Home	Sports
Gender						
Female		4994	5015	4984	5017	5104
Male		4979	4954	4963	4983	5007

Chi-Square Statistic: 0.4467  
Degrees of Freedom: 4  
P-Value: 0.9785  
Accept Null Hypothesis: No significant relationship between Gender and Product Category.

```
[ ]: df.columns
```

```
[ ]: Index(['Order_ID', 'Customer_ID', 'Gender', 'Age', 'Country',
       'Product_Category', 'Quantity', 'Unit_Price', 'Discount', 'Order_Date',
       'Payment_Method', 'Ad_Campaign', 'Returned', 'Total_Sales', 'Year',
       'Month_Name', 'Cost', 'Profit', 'Unit_Cost', 'Month_Year'],
      dtype='object')
```

```
[ ]: df[["Gender", "Product_Category"]].value_counts().unstack()
```

	Product_Category	Electronics	Fashion	Groceries	Home	Sports
Gender						
Female		4994	5015	4984	5017	5104
Male		4979	4954	4963	4983	5007

### 0.1.3 Business Insights & Recommendations

#### Top Products (Revenue & Profit)

- Electronics (20.2%), Sports (20.1%), and Fashion (20.0%) together drive ~60% of total revenue.
- **Recommendation:** Prioritize inventory allocation, promotional spend, and marketing focus on these categories.

#### Underperforming Categories

- Home (19.8%) and Groceries (19.9%) contribute the least to overall revenue and profit.
- **Recommendation:** Run targeted pricing/discount campaigns; if performance doesn't improve, consider rationalizing SKUs.

#### Customer Gender Insights

- Female customers generated **\$594.9M**, slightly higher than males at **\$584.6M**.
- **Recommendation:** Develop gender-specific campaigns — loyalty programs for females, product bundling for males.

#### Seasonality Trends

- Sales peak during **Nov–Dec** (festive season) and dip in **Mar–Apr**.
- **Recommendation:** Increase stock levels and intensify promotions in festive months; optimize costs in low-demand periods.

#### High Volume High Profitability

- Example: Groceries record high sales volume but deliver the lowest margin (~19.8% of profit).
- **Recommendation:** Improve margins through pricing adjustments, product bundling, or promoting premium/high-value alternatives.